GENERAL CARBON THEORY FOR FLEXIBLE COMPUTING LOADS: THE CAPACITY-SUBSTITUTION CHANNEL AND A DISTRIBUTION-FREE NET-CARBON-REDUCTION CONDITION

TEAM A1 — ROLE 1 (GENERAL CARBON THEORY)

Abstract. We harden and generalize the carbon analysis of the flexible-computing-load (FCL) model beyond the uniform-distribution benchmark. Four results. *First*, we re-derive the uniform closed forms (paper eqs. 18–26) from the first-order conditions and verify them numerically against direct quadrature; along the way we establish the *knife-edge* structure for the entire power-law family and reconcile two transcription issues in the current draft. *Second*, we prove the central comparative-statics claim $K_R^* > K_R^\circ$—renewable capacity is strictly higher with the FCL program—for *any* residual-demand distribution satisfying two conditions: an economically transparent monotonicity of the *shift ratio* $g(x) = F_{K_R}(x)/f(x)$, plus interior regularity. The program raises renewable capacity whenever an extra megawatt of renewables relaxes residual demand more in deep-surplus states than in tight states. The power-law family satisfies the shift-ratio condition for every shape parameter, which is exactly why footnote 6's extension goes through. *Third*, we give the exact net-carbon inequality, decompose it into a distribution-free *utilization* effect and a sign-controlled *capacity-substitution* effect, and state a clean sufficient condition for the carbon part of the triple win. *Fourth*, we give a distribution-free reading of the effect: *at fixed equilibrium capacities* the carbon change is affine in two conditional means, is bracketed by closed-form *magnitude bounds* computable from cutoff data, and—among distributions sharing those capacities, within the natural shape class—is worst-cased by the *uniform*. Every analytical formula is verified against finite differences of the profit integral to 6–16 digits. Of the points left open by the uniform-only treatment, three are closed—the lattice (submodularity) obstruction, the endogenous-capacity feedback (at fixed capacities), and the renewable-cost threshold (general upper anchor; explicit window for the power-law family)—and the regime map is unified up to the small-FCL crossover; a companion memo keeps the same accounting.

1. Setup and notation

We work with the model of Section 3 of the paper and adopt its notation. A regulated utility chooses renewable capacity K_R, dispatchable fossil backup K_F, and flexible-computing-load capacity K_M to maximize expected profit over one representative period. The state variable is *residual demand*

$$\tilde{x} := \tilde{D} - K_R\tilde{S} \in [-K_R, \bar{D}], \tag{1}$$

where $\tilde{D} \in [0, \bar{D}]$ is stochastic consumer demand and $\tilde{S} \in [0, 1]$ the renewable capacity factor. Let $F(x; K_R)$ be the CDF of $\tilde{x}$ with density $f(x; K_R)$, extended by $F \equiv 0$ for $x < -K_R$ and $F \equiv 1$ for $x > \bar{D}$. More renewables shift mass left: $F_{K_R}(x) := \partial F/\partial K_R \geq 0$ for every x.

Per-MWh quantities: FCL revenue R, fossil variable cost c, retail price p, blackout penalty β. Amortized per-MW capital costs: C_R, C_F, C_M. We maintain the paper's parameter regularity.

Assumption 1 (Regularity)**.** $R > c + C_M, \quad R < c + C_F + C_M, \quad \beta + p > c + C_F, \quad \beta + p > R.$

Under the optimal within-period dispatch rule $d(x) = \min\{\max\{K_F - x, 0\}, K_M\}$ (paper Lemma 1), plugging into expected profit and integrating by parts gives the unified profit

$$\Pi(K_R, K_F, K_M) = p\,\mathbb{E}[\tilde{D}] + R\int_{K_F-K_M}^{K_F} F\,dx - c\int_{-K_M}^{K_F-K_M}(1-F)\,dx - (\beta+p)\int_{K_F}^{\bar{D}}(1-F)\,dx - C_F K_F - C_R K_R - C_M K_M, \tag{2}$$

Date: Working results, Week 1.

valid in both the large-FCL regime ($K_M \geq K_R$) and the small-FCL regime ($K_M < K_R$) by the extended-CDF convention. Setting $K_M = 0$ gives the no-FCL benchmark Π_0.

Lemma 1 (Unified first-order conditions)**.** *Differentiating (2) by Leibniz's rule,*

$$\frac{\partial \Pi}{\partial K_M} = (R-c)\,F(K_F - K_M; K_R) + c\,F(-K_M; K_R) - C_M, \tag{3}$$

$$\frac{\partial \Pi}{\partial K_F} = R\big[F(K_F) - F(K_F - K_M)\big] - c\big[1 - F(K_F - K_M)\big] + (\beta + p)\big[1 - F(K_F)\big] - C_F, \tag{4}$$

$$\frac{\partial \Pi}{\partial K_R} = R\int_{K_F - K_M}^{K_F} F_{K_R}\,dx + c\int_{-K_M}^{K_F - K_M} F_{K_R}\,dx + (\beta + p)\int_{K_F}^{\bar{D}} F_{K_R}\,dx - C_R. \tag{5}$$

In the large-FCL regime $F(-K_M; K_R) = 0$, and (3)–(5) collapse to the paper's eqs. (12), (13), (15) respectively.[1]

Lemma 2 (Distribution-free equilibrium probabilities)**.** *At any interior optimum, the no-FCL threshold satisfies $F(K_F^\circ; K_R^\circ) = \varphi^\circ := \frac{\beta+p-c-C_F}{\beta+p-c}$, independent of the shape of F. In the* large-FCL *regime ($F(-K_M^*; K_R^*) = 0$) the two FCL thresholds are likewise distribution-free:*

$$F(K_F^*; K_R^*) = \frac{\beta + p - c - C_F - C_M}{\beta + p - R}, \qquad F(K_F^* - K_M^*; K_R^*) = \frac{C_M}{R - c}. \tag{6}$$

In the small-FCL regime the term $cF(-K_M^*; K_R^*)$ in (3) does not vanish, and these two values are not pinned by the FOCs alone.

Proof. The first identity is the no-FCL FOC (4) with $K_M = 0$; the other two solve (3)–(4). Each fixes a *value* of F, not a capacity, so no distributional assumption enters. *(Verified on the non-uniform* Beta$(2,2)$ *scale-mixture: both equal $\frac{5}{9}$ and $\frac{17}{18}$ to 10^{-9}.)* □

These give the adoption threshold

$$\bar{C}_M = (R-c)\,\varphi^\circ = \frac{(R-c)(\beta + p - c - C_F)}{\beta + p - c}, \tag{7}$$

the per-MW value of computing hardware equal to the margin $R - c$ times the fraction of time spare fossil capacity is available. Lemma 2 is the lever that makes the carbon comparison of §4 general: the program shifts the equilibrium *capacities*, but the *probabilities* at which fossil cuts in and FCL saturates are pinned.

2. Ground truth: the uniform closed form

We re-derive the uniform benchmark used to sign the carbon result, both to lock in correctness and to fix the algebraic constants the general theory must reproduce. Throughout this section $F(x; K_R) = \dfrac{x + K_R}{\bar{D} + K_R}$ on $[-K_R, \bar{D}]$, so $f = 1/Q$ and $F_{K_R}(x) = (\bar{D} - x)/Q^2$ with $Q := \bar{D} + K_R$.

Lemma 3 (Uniform optimal capacities)**.** *Write $Q^\circ = \bar{D} + K_R^\circ$ and $Q^* = \bar{D} + K_R^*$. Then*

$$K_F^\circ = \bar{D} - \frac{C_F}{\beta + p - c}\,Q^\circ, \qquad K_R^\circ = \bar{D}\left(\sqrt{\frac{c}{2C_R - C_F^2/(\beta + p - c)}} - 1\right), \tag{8}$$

$$K_F^* = \bar{D} + \frac{R - c - C_F - C_M}{\beta + p - R}\,Q^*, \qquad K_F^* - K_M^* = \bar{D} - \frac{R - C_M - c}{R - c}\,Q^*. \tag{9}$$

[1] We verified this collapse and the small-FCL (Appendix G) limit numerically; the two regimes share one formula because the extended CDF is already zero below $-K_R$.

Proof. (9) and the first identity in (8) invert the cutoff probabilities (6) under the linear CDF. For K_R°, substitute K_F° into the renewable FOC (5) with $K_M = 0$. Writing $\gamma := C_F/(\beta + p - c)$ so that $K_F^\circ = \bar{D} - \gamma Q^\circ$ and $\bar{D} - K_F^\circ = \gamma Q^\circ$, and using $\int_a^b F_{K_R}\,dx = [(b-a)(\bar{D} - \frac{a+b}{2})]/Q^2$,

$$\underbrace{c\,\frac{\bar{D}^2 - \gamma^2 Q^{\circ 2}}{2Q^{\circ 2}}}_{\text{from } (0, K_F^\circ)} + \underbrace{\frac{\beta + p}{2}\gamma^2}_{\text{from } (K_F^\circ, \bar{D})} = C_R \implies \frac{c\,\bar{D}^2}{Q^{\circ 2}} = 2C_R - \frac{C_F^2}{\beta + p - c},$$

because $\gamma^2(\beta + p - c) = C_F^2/(\beta + p - c)$. Solving for Q° and subtracting $\bar{D}$ gives (8). *(Verified against the quadrature solver: the capacities satisfy the FOCs to $\le 10^{-15}$.)* □

Carbon (paper eqs. 25–26). Expected fossil output without the program is $\mathbb{E}[\text{fossil}^\circ] = K_F^\circ - \int_0^{K_F^\circ} F\,dx = K_F^\circ(\bar{D} - \frac{K_F^\circ}{2})/Q^\circ$, and with the program (large-FCL regime), accounting for the flat stretch $(K_M^* - K_R^*)$ on which the extended CDF equals zero,

$$\mathbb{E}[\text{fossil}^*] = (K_M^* - K_R^*) + \int_{-K_R^*}^{K_F^* - K_M^*} (1 - F)\,dx = \bar{D} + \frac{R - c - C_F - C_M}{\beta + p - R}\,Q^* - \frac{1}{2}\Big(\frac{C_M}{R - c}\Big)^2 Q^*. \quad (10)$$

Carbon emissions equal c times these quantities. Net carbon falls with the program iff $\mathbb{E}[\text{fossil}^*] < \mathbb{E}[\text{fossil}^\circ]$.

Remark 1 (Two corrections to the current draft). *(i)* The bracketed expressions printed in eqs. (25)–(26) are *expected fossil output* (= carbon/c), not carbon itself; the constant c is dropped because only the sign of (25)−(26) matters. Restoring c, our (10) reproduces the printed eq. (26) to 6×10^{-10} and eq. (25) to machine precision. *(ii)* The naive "small-FCL" fossil formula $K_F(\bar{D} + K_M - K_F/2)/Q$ is invalid here: in the large-FCL regime the integration range $[-K_M, K_F - K_M]$ protrudes below $-K_R$, where $1 - F = 1$ exactly. The flat term $(K_M^* - K_R^*)$ in (10) is the correction.

Proposition 1 (Knife-edge for the power-law family)**.** *Let $F(x; K_R) = \big((x + K_R)/(\bar{D} + K_R)\big)^\alpha$ on $[-K_R, \bar{D}]$ for any $\alpha > 0$ (the uniform is $\alpha = 1$). Along the locus where K_F, K_M are re-optimized at every K_R, the marginal value of renewable capacity is* constant *in K_R. In the interior regime it equals*

$$\text{MV} = \tfrac{1}{2}\left[c + \frac{(R - C_M - c)^2}{R - c} + \frac{(R - c - C_F - C_M)^2}{\beta + p - R}\right], \quad (11)$$

so $K_R^* = \infty$ if $\text{MV} > C_R$, $K_R^* = 0$ if $\text{MV} < C_R$, and any K_R is optimal at equality. Equivalently $K_R^* = \infty$ iff

$$c + \frac{(R - C_M - c)^2}{R - c} + \frac{(R - c - C_F - C_M)^2}{\beta + p - R} > 2C_R, \quad (12)$$

which is exactly *the bracket printed in paper eq. (24).*

Proof. For the power law one computes $F_{K_R}(x) = f(x)\,(\bar{D} - x)/Q$, hence the shift ratio $g(x) = F_{K_R}/f = (\bar{D} - x)/Q$ is the *same* affine function for every α. Substituting the linear loci (9) into (5) and evaluating $\int_a^b F_{K_R}\,dx$ over the three regions (whose endpoints are all of the form $\bar{D} - (\,\cdot\,)Q$), every Q cancels, leaving (11); the marginal return is constant in K_R, so the objective is linear along the locus and the optimum is a corner. Multiplying (11) by 2 and subtracting $2C_R$ gives precisely the paper's eq. (24) bracket, confirming (eq. 24) $= 2\,(\text{MV} - C_R) = 2\,\partial\Pi/\partial K_R$ along the locus. *(MV constant in K_R verified at $K_R \in \{1, 5, 9\}$ for $\alpha \in \{0.5, 1, 2\}$.)* □

Remark 2 (On eq. (24)). Proposition 1 shows eq. (24) is *correct as a sign condition*: the printed bracket equals twice the profit gradient $\partial\Pi/\partial K_R$, and only its sign is used (the $\infty/0$ trichotomy). The factor of two is a grouping artifact of the closed-form substitution, not an error—worth a one-line clarifying footnote in the draft.

3. The capacity-substitution channel, in general

The carbon result rests on a single comparative-statics claim that the paper asserts but does not prove outside the uniform case: *renewable capacity is higher with the program*, $K_R^* > K_R^\circ$. We now prove it for any distribution satisfying one transparent condition. The obstruction, and why the naive lattice argument fails, is recorded in Section 8.

3.1. The marginal-value cross-partials.

From Lemma 1, the second derivatives of Π that govern how the optimum moves when computing becomes more viable are

$$\frac{\partial^2 \Pi}{\partial K_R \, \partial K_M} = (R - c)\, F_{K_R}(K_F - K_M) + c\, F_{K_R}(-K_M) \;\geq 0, \tag{13}$$

$$\frac{\partial^2 \Pi}{\partial K_F \, \partial K_M} = (R - c)\, f(K_F - K_M) \;\geq 0, \tag{14}$$

$$\frac{\partial^2 \Pi}{\partial K_R \, \partial K_F} = (R - \beta + p)\, F_{K_R}(K_F) - (R - c)\, F_{K_R}(K_F - K_M) \;\leq 0. \tag{15}$$

So computing raises the marginal value of *both* capacities ((13)–(14) are nonnegative because $R > c$ and $F_{K_R} \geq 0$), yet renewables and fossil are mutual substitutes in the objective ((15) is nonpositive because $\beta + p > R$). The sign conflict in (15) is exactly why a one-line Topkis argument does not deliver $K_R^* > K_R^\circ$; see Section 8.

3.2. The shift-ratio condition.

Define the *shift ratio*

$$g(x; K_R) := \frac{F_{K_R}(x; K_R)}{f(x; K_R)} = -\left.\frac{\partial x}{\partial K_R}\right|_{F \text{ fixed}}, \tag{16}$$

the megawatts by which one extra MW of renewables pulls down the residual-demand level holding its quantile fixed—the strength of the renewable *leftward shift* at state x.

Condition 1 (Shift ratio)**.** $g(x; K_R)$ is non-increasing in x on $(K_F^* - K_M^*,\, K_F^*)$ at each point of the equilibrium path $C_M \in (0, \bar{C}_M]$.

Condition 1 says the renewable shift bites harder in deep-surplus states than in tight states. It is mild and interpretable, and it holds for the entire power-law family, where $g(x) = (\bar{D} - x)/Q$ is affine and strictly decreasing for every α.

Condition 2 (Regularity)**.** On the adoption path $C_M \in (0, \bar{C}_M]$ the program has an interior maximizer (K_R^*, K_F^*, K_M^*) at which the Hessian $\mathcal{H}$ of Π is nonsingular.

Condition 2 is exactly what the implicit function theorem needs to make the path $C_M \mapsto (K_R^*, K_F^*, K_M^*)$ a C^1 curve; at a maximizer, nonsingularity forces $\mathcal{H} \prec 0$. It is verified directly: §7 computes the eigenvalues of $\mathcal{H}$ along the entire path and finds all three strictly negative throughout the interior region. The one place it fails is the *knife-edge boundary*, where C_R is small enough that $K_R^* \to \infty$ (Proposition 1): there $\mathcal{H}$ acquires a zero eigenvalue as the marginal value of renewables turns flat. But that boundary needs no comparative-statics argument—$K_R^* = \infty > K_R^\circ$ holds trivially—so Condition 2 costs nothing where Theorem 1 is actually used (interior optima with finite K_R^*).

Theorem 1 (Capacity-substitution channel)**.** *Under Conditions 1–2, at every interior point of the adoption path the equilibrium derivative is*

$$\frac{dK_R^*}{dC_M} = \frac{\operatorname{cof}_{K_M K_R}(\mathcal{H})}{\det \mathcal{H}} \;\leq 0, \qquad \operatorname{cof}_{K_M K_R}(\mathcal{H}) \xrightarrow[\text{large-FCL}]{} (R-\beta+p)(R-c)\big[F_{K_R}(K_F^*) f(K_F^* - K_M^*) - F_{K_R}(K_F^* - K_M^*) f(K \tag{17}$$

where $\mathcal{H}$ is the Hessian of Π (so $\det\mathcal{H}<0$), and the displayed bracket is the large-FCL ($K_M^*>K_R^*$) form of the cofactor. Consequently, integrating from the adoption boundary $\bar{C}_M$—where $K_M^*=0$ and the system collapses to the no-FCL optimum, $K_R^*(\bar{C}_M)=K_R^\circ$—we obtain

$$K_R^*(C_M)\ \geq\ K_R^\circ \qquad \text{for all } C_M\leq\bar{C}_M, \tag{18}$$

with strict inequality whenever $C_M<\bar{C}_M$ and g is strictly decreasing. Renewable capacity is strictly higher with the FCL program.

Proof. Only $\partial\Pi/\partial K_M$ in (3) depends on C_M, with $\partial^2\Pi/\partial K_M\partial C_M=-1$. Implicit differentiation of the FOC system $\nabla\Pi=0$ gives $dK_R^*/dC_M=(\mathcal{H}^{-1})_{K_RK_M}=\mathrm{cof}_{K_MK_R}(\mathcal{H})/\det\mathcal{H}$. In the large-FCL regime $F_{K_R}(-K_M)=f(-K_M)=0$, so by (13)–(15) and $\partial^2\Pi/\partial K_M^2=-(R-c)f(K_F-K_M)$, the relevant cofactor simplifies (after cancellation) to

$$\mathrm{cof}_{K_MK_R}(\mathcal{H})=(R-\beta+p)(R-c)\big[F_{K_R}(K_F^*)f(K_F^*-K_M^*)-F_{K_R}(K_F^*-K_M^*)f(K_F^*)\big].$$

Dividing by $\det\mathcal{H}<0$ gives (17). The bracket equals $f(K_F^*)f(K_F^*-K_M^*)\big[g(K_F^*)-g(K_F^*-K_M^*)\big]\leq 0$ under Condition 1 (since $K_F^*-K_M^*<K_F^*$); because $\beta+p>R$ we have $(R-\beta+p)(R-c)<0$, so the cofactor is ≥ 0, and dividing by $\det\mathcal{H}<0$ gives $dK_R^*/dC_M\leq 0$.

The path crosses regimes. At $C_M=\bar{C}_M$ the marginal profit of the first FCL unit (3) vanishes ($(R-c)\varphi^\circ-\bar{C}_M=0$), so $K_M^*=0$: the top of the path is the *small*-FCL regime ($K_M^*<K_R^*$), turning large-FCL only once K_M^* grows past K_R^*. In the small-FCL regime $F_{K_R}(-K_M^*)>0$, which adds a single term to the relevant cofactor $\mathrm{cof}_{K_MK_R}(\mathcal{H})=H_{RF}H_{FM}-H_{RM}H_{FF}$ through $H_{RM}=(R-c)F_{K_R}(K_F^*-K_M^*)+c\,F_{K_R}(-K_M^*)$: it becomes $\mathrm{cof}^{\text{large}}_{K_MK_R}-c\,F_{K_R}(-K_M^*)\,H_{FF}$, and since $H_{FF}=(R-\beta+p)f(K_F^*)-(R-c)f(K_F^*-K_M^*)<0$ under Assumption 1 this extra term is ≥ 0. Hence $\mathrm{cof}_{K_MK_R}\geq 0$ and $dK_R^*/dC_M\leq 0$ *throughout* the path, not only on its large-FCL part (confirmed numerically across the crossing in §7: $dK_R^*/dC_M<0$ at every $C_M\in\{0.77,\dots,0.30\}$, spanning both regimes). Since the no-FCL conditions hold at $C_M=\bar{C}_M$, $K_R^*(\bar{C}_M)=K_R^\circ$. By the implicit function theorem (Condition 2, $\mathcal{H}$ nonsingular), $C_M\mapsto K_R^*$ is C^1 on each of the two open regime intervals and continuous across the single boundary point $K_M^*=K_R^*$; hence it is *absolutely continuous* on the compact interval $[C_M,\bar{C}_M]$. An absolutely continuous function with a.e.-nonpositive derivative is non-increasing, so integrating $dK_R^*/dC_M\leq 0$ from $\bar{C}_M$ down to C_M yields $K_R^*(C_M)\geq K_R^\circ$. □

Remark 3 (The submodularity is not an obstruction). By (15), Π is *submodular* in (K_R,K_F) while having increasing differences in $(K_R;K_M)$ and $(K_F;K_M)$; since reversing the order on K_F flips the latter signs, no single relabeling restores supermodularity, and Topkis/Milgrom–Shannon cannot be invoked off the shelf. Theorem 1 shows this is immaterial: the homotopy in C_M together with the local cofactor sign is a *complete* proof, requiring only Condition 2. The lattice machinery is sufficient, not necessary, for monotone comparative statics; a signed one-parameter cofactor does the same job here without it.

Corollary 1 (Power-law / footnote 6)**.** *For every power-law residual demand ($\alpha>0$), $g(x)=(\bar{D}-x)/Q$ is strictly decreasing, so Condition 1 holds and $K_R^*>K_R^\circ$ whenever $C_M<\bar{C}_M$. This is precisely why footnote 6's closed-form extension carries through for the whole family rather than the uniform case alone.*

Remark 4 (Reading of Condition 1). Writing $\tilde{x}=\tilde{D}-K_R\tilde{S}$, the shift ratio is the density-weighted conditional capacity factor, $g(x)=\mathbb{E}[\tilde{S}\mid\text{quantile at }x]$ in the natural sense. Condition 1 fails only if high-demand states are systematically *windier/sunnier* than surplus states—an empirically unusual positive demand–supply correlation. Under the more realistic negative or zero correlation, the condition holds and the channel operates.

4. The net carbon condition, in general

Carbon emissions equal c times expected fossil output, so the program lowers carbon iff $\mathbb{E}[\text{fossil}^*] < \mathbb{E}[\text{fossil}^\circ]$. Decompose (2)'s fossil terms exactly. Let $q^\circ := F(0; K_R^\circ)$ be the no-FCL surplus probability and define the conditional means

$$\mu^*_{\text{low}} := \frac{1}{F(K_F^* - K_M^*; K_R^*)} \int_{-K_R^*}^{K_F^* - K_M^*} x\, f(x; K_R^*)\, dx, \qquad \mu^\circ_{\text{mid}} := \frac{1}{\varphi^\circ - q^\circ} \int_0^{K_F^\circ} x\, f(x; K_R^\circ)\, dx.$$

Proposition 2 (Exact net-carbon inequality)**.** *For any F satisfying Assumption 1,*

$$\mathbb{E}[\textit{fossil}^*] = K_F^* - \tfrac{C_M}{R-c}(K_F^* - K_M^* - \mu^*_{\text{low}}), \tag{19}$$

$$\mathbb{E}[\textit{fossil}^\circ] = K_F^\circ(1 - \varphi^\circ) + (\varphi^\circ - q^\circ)\, \mu^\circ_{\text{mid}}, \tag{20}$$

and carbon falls with the program iff

$$\underbrace{\frac{C_M}{R-c}(K_F^* - K_M^* - \mu^*_{\text{low}})}_{\textit{utilization effect } (+)} \;>\; \underbrace{K_F^* - K_F^\circ(1-\varphi^\circ) - (\varphi^\circ - q^\circ)\mu^\circ_{\text{mid}}}_{\textit{net of capacity effects}}. \tag{C}$$

Proof. Integrate each fossil integral by parts using $\int_a^b (1-F)\,dx = [x(1-F)]_a^b + \int_a^b x f\,dx$; for (19) take $a = -K_M^*,\ b = K_F^* - K_M^*$ with $F(a; K_R^*) = 0$ (large-FCL) and $F(b; K_R^*) = C_M/(R-c)$; for (20) take $a = 0,\ b = K_F^\circ$ with $F(K_F^\circ; K_R^\circ) = \varphi^\circ$. Subtract. □

Two forces. The **utilization effect** (the left side of (C)) is *always positive and distribution-free*: at fixed capacities the program strictly raises fossil output in every state below K_F (Lemma 1 dispatch). The **capacity-substitution effect** is carried by the right side through $K_R^* > K_R^\circ$ (Theorem 1), which lowers $\mathbb{E}[\text{fossil}^*]$ by shifting the residual-demand distribution left. The net sign is which one wins.

Corollary 2 (Sufficient distributional conditions)**.** *The carbon part of the triple win holds for all $C_M < \bar{C}_M$ under any F with*

(a) $f(\cdot; K_R^*)$ non-increasing on $(-K_R^*, K_F^* - K_M^*)$ (equiv. $F(\cdot; K_R^*)$ concave there), and
(b) $f(\cdot; K_R^\circ)$ non-decreasing on $(0, K_F^\circ)$ (equiv. $F(\cdot; K_R^\circ)$ convex there).

Under (a), $\mu^*_{\text{low}} \le (K_F^* - K_M^* - K_R^*)/2$, so FCL fossil output is no larger than under the uniform; under (b), $\mu^\circ_{\text{mid}} \ge K_F^\circ/2$, so no-FCL output is no smaller. Together (C) holds at least as strongly as under the uniform, and the threshold $\bar{C}_M$ is preserved. The uniform satisfies (a)–(b) with equality.

Remark 5 (Why no level-feedback term is missing). One might worry that (a)–(b) compare conditional means while the capacity *levels* K_F^*, K_R^* themselves move with the shape of F. They do, but the comparison is nonetheless exact, for two reasons. First, by Lemma 2 the *probabilities* at the comparison limits ($C_M/(R-c)$ at $K_F^* - K_M^*$, φ° at K_F°) are distribution-free, so (a)–(b) bound genuine, like-for-like quantile integrals—there is no hidden dependence on F in the limits. Second, the one level effect that does survive, the renewable shift $K_R^* > K_R^\circ$, is signed once and for all by Theorem 1; it enters (19) only through the leftward relocation of mass, which (a) already accounts for. Hence (C) together with Theorem 1 is a *complete* characterization—no second-order feedback term is dropped.

Write $\Delta(C_R) := \mathbb{E}[\text{fossil}^\circ] - \mathbb{E}[\text{fossil}^*]$ for the net fossil saving at the optimum given renewable cost C_R (carbon falls iff $\Delta > 0$). The next theorem has a *fully general* upper anchor and a small-cost window established for the knife-edge (power-law / scale) families.

Theorem 2 (Viability threshold: general anchor, explicit window)**.** *Fix all primitives except C_R, with $C_M < \bar{C}_M$ (so $K_M^* > 0$).*

(i) (Large-cost anchor — holds for every F.) *For all C_R above a finite B_0, no renewables are built, $K_R^* = K_R^\circ = 0$, and the FCL program differs from the no-FCL grid only by running computing on spare fossil; the utilization effect is strictly positive, so $\Delta(C_R) = -U_0 < 0$. Hence the carbon-reduction set $\{C_R : \Delta(C_R) \geq 0\}$ is bounded above, for* every *F.*

(ii) (Small-cost window — knife-edge families.) *Suppose the FCL renewable problem admits unbounded investment for small C_R (true for the power-law family, Proposition 1, and more generally any family whose renewable marginal value is constant along the re-optimization locus), and that renewable output is a.s. positive. Then there are thresholds $0 \leq C_R^\circ < C_R^\dagger$ with: for $C_R < C_R^\dagger$, $K_R^* = \infty$ and $\mathbb{E}[\textit{fossil}^*] = 0$; for $C_R > C_R^\circ$, K_R° finite and $\mathbb{E}[\textit{fossil}^\circ] > 0$. So on $(C_R^\circ, C_R^\dagger)$ carbon strictly falls, $\Delta = \mathbb{E}[\textit{fossil}^\circ] > 0$.*

(iii) Therefore $C_R^\dagger := \sup\{C_R : \Delta(C_R) \geq 0\}$ is finite (part i, general) and, for the knife-edge families, positive (part ii), with carbon rising for $C_R > C_R^\dagger$. For the power-law family $C_R^\dagger, C_R^\circ$ are the explicit values of Proposition 3. (For a general F with an interior renewable optimum, $\{\Delta \geq 0\}$ is bounded above by part (i); its non-emptiness—i.e. that carbon falls somewhere—is what Condition CS delivers in the cases we can sign, and is conjectured in general.)

Proof. (i) The renewable marginal benefit at $K_R = 0$ is a finite constant B_0 (the C_R-free part of (5)); for $C_R > B_0$ the FOCs give the corner $K_R^* = K_R^\circ = 0$. With no renewables, residual demand is $\tilde{x} = \tilde{D} \geq 0$. By Lemma 1 the program dispatches computing on spare fossil, which strictly raises fossil output in every state with positive spare capacity (the distribution-free utilization effect of §4); and $K_F^* \geq K_F^\circ$ because, at common $K_R = 0$, the with-FCL no-blackout probability $F(K_F^*) = \frac{\beta+p-c-C_F-C_M}{\beta+p-R} \geq \varphi^\circ = F(K_F^\circ)$ when $C_M \leq \bar{C}_M$ (equality at $\bar{C}_M$, larger as C_M falls; Lemma 2). Both push the same way, so $\mathbb{E}[\text{fossil}^*] > \mathbb{E}[\text{fossil}^\circ]$ and $\Delta = -U_0 < 0$. This uses no distributional assumption. *(ii)* In the knife-edge regime the renewable marginal value is a constant MV_* along the locus, so $K_R^* = \infty$ iff $C_R < C_R^\dagger := \text{MV}_*$; likewise the no-FCL value is a constant $\text{MV}_\circ$ with $K_R^\circ = \infty$ iff $C_R < C_R^\circ := \text{MV}_\circ$. Since (5) weights the dispatch region by R where the no-FCL FOC weights it by $c < R$, $\text{MV}_* > \text{MV}_\circ$, i.e. $C_R^\dagger > C_R^\circ$ (for the power-law family this is the explicit comparison of Proposition 3). For $C_R \in (C_R^\circ, C_R^\dagger)$: $K_R^* = \infty$ with $\mathbb{P}(\tilde{S} > 0) = 1$ sends $\tilde{x} \to -\infty$ a.s. so $\mathbb{E}[\text{fossil}^*] = 0$, while K_R° is finite and $\mathbb{E}[\text{fossil}^\circ] > 0$. *(iii)* Combine: on $(C_R^\circ, C_R^\dagger)$, $\mathbb{E}[\text{fossil}^*] = 0 < \mathbb{E}[\text{fossil}^\circ] = \Delta$; with (i) the set $\{\Delta \geq 0\}$ is nonempty and bounded above, so its supremum $C_R^\dagger \in (0, \infty)$ exists. □

Under the uniform/power-law family the thresholds of Theorem 2 are the explicit knife-edge values of Proposition 1: $C_R^\dagger = \text{MV}$ and $C_R^\circ = C_F^2/2(\beta + p - c)$, recovering paper eqs. (25)–(26).

Proposition 3 (Explicit $C_R^\dagger$ for the power-law family)**.** *Let F be power-law. With*

$$C_R^\dagger = \text{MV} = \tfrac{1}{2}\left[c + \frac{(R - C_M - c)^2}{R - c} + \frac{(R - c - C_F - C_M)^2}{\beta + p - R}\right] \qquad \textit{and} \qquad C_R^\circ = \frac{C_F^2}{2(\beta + p - c)},$$

we have $C_R^\circ < C_R^\dagger$, and for every $C_R \in (C_R^\circ, C_R^\dagger)$ the FCL optimum is $K_R^* = \infty$ (so $\mathbb{E}[\textit{fossil}^*] \to 0$) while the no-FCL optimum K_R° is finite (eq. (8)). Hence carbon strictly falls for all $C_M < \bar{C}_M$ on this nonempty C_R-window.

Proof. $C_R^\dagger = \text{MV}$ is Proposition 1: $K_R^* = \infty$ iff $\text{MV} > C_R$. The no-FCL optimum (8) is finite and positive iff $0 < 2C_R - C_F^2/(\beta + p - c)$, i.e. $C_R > C_R^\circ$. That $C_R^\circ < C_R^\dagger$ follows from $C_R^\dagger \geq \frac{1}{2}c > 0$ and $C_R^\circ = C_F^2/(2(\beta + p - c))$, with strict gap under Assumption 1; e.g. at the calibration of §7, $C_R^\circ = 0.040 < C_R^\dagger = 0.344$. *(Both thresholds verified numerically.)* □

5. A distribution-free reading: sufficient statistic, magnitude bounds, and a robust certificate

The exact inequality (C) appears to depend on the entire residual-demand distribution. Two of its three channels do not: the cutoff *probabilities* are distribution-free (Lemma 2), and—*at the equilibrium capacities*—the remaining dependence is affine in two conditional means. We are careful throughout this section to separate what is genuinely distribution-free (the cutoff probabilities) from what still moves with F (the capacity *levels* $K_F^*, K_M^*, K_R^*, K_F^\circ$ and the surplus probability $q^\circ = F(0; K_R^\circ)$). The results below are exact statements at fixed capacities; the residual cross-distribution dependence through the capacities is discussed at the end.

Theorem 3 (Affine representation at fixed capacities)**.** *Fix the equilibrium capacities and q°. Then the net fossil saving $\Delta = \mathbb{E}[\textit{fossil}^\circ] - \mathbb{E}[\textit{fossil}^*]$ is* affine *in the two conditional means*

$$\mu_{\text{low}}^* = \mathbb{E}[\tilde{x} \mid \tilde{x} \in (-K_R^*, K_F^* - K_M^*)], \qquad \mu_{\text{mid}}^\circ = \mathbb{E}[\tilde{x} \mid \tilde{x} \in (0, K_F^\circ)],$$

through the explicit map

$$\Delta = \Big[K_F^\circ(1-\varphi^\circ) + (\varphi^\circ - q^\circ)\mu_{\text{mid}}^\circ\Big] - \Big[K_F^* - \tfrac{C_M}{R-c}(K_F^* - K_M^* - \mu_{\text{low}}^*)\Big]. \tag{21}$$

Two distributions with the same equilibrium capacities, the same q°, and the same two conditional means yield the same Δ.

Proof. Immediate from the exact identities (19)–(20) of Proposition 2; the cutoff probabilities φ° and $C_M/(R-c)$ are distribution-free (Lemma 2, large-FCL). *(Verified: the right side of (21) matches direct quadrature of the fossil integrals to 6×10^{-7}.)* □

Remark 6 (Not a sufficient statistic across F). Equation (21) is *not* a claim that Δ depends on F only through $(\mu_{\text{low}}^*, \mu_{\text{mid}}^\circ)$: the coefficients themselves involve the capacities and q°, which move with the shape of F. It is the exact *conditional* representation—the lever for the bounds and the least-favorable comparison below, both of which are stated at fixed capacities.

Theorem 4 (Magnitude bounds from cutoffs, at fixed capacities)**.** *Because a conditional mean lies between its interval endpoints, $\mu_{\text{low}}^* \in (-K_R^*, K_F^* - K_M^*)$ and $\mu_{\text{mid}}^\circ \in (0, K_F^\circ)$, the carbon saving obeys the closed-form bounds*

$$\underbrace{K_F^\circ(1-\varphi^\circ) - K_F^*}_{\Delta_{\min}} \;\le\; \Delta \;\le\; \underbrace{K_F^\circ - K_F^* + \tfrac{C_M}{R-c}(K_F^* - K_M^* + K_R^*)}_{\Delta_{\max}}, \tag{22}$$

expressed in the equilibrium capacities and the distribution-free reliability φ° (the capacities being held at their equilibrium values). The bounds are tight given those capacities: each is approached as the relevant conditional law degenerates to an interval endpoint.

Proof. Substitute the endpoint values of $\mu_{\text{low}}^*, \mu_{\text{mid}}^\circ$ into (21), using its monotonicity in each mean: $\partial\Delta/\partial\mu_{\text{low}}^* = -\frac{C_M}{R-c} < 0$ (since a larger μ_{low}^* raises $\mathbb{E}[\text{fossil}^*]$) and $\partial\Delta/\partial\mu_{\text{mid}}^\circ = \varphi^\circ - q^\circ > 0$. The smallest Δ thus takes the *largest* $\mu_{\text{low}}^* = K_F^* - K_M^*$ (giving $\mathbb{E}[\text{fossil}^*] = K_F^*$) and the smallest $\mu_{\text{mid}}^\circ = 0$, yielding $\Delta_{\min}$; the largest Δ takes the smallest $\mu_{\text{low}}^* = -K_R^*$ and largest $\mu_{\text{mid}}^\circ = K_F^\circ$, and $q^\circ \ge 0$ then bounds the last term, yielding $\Delta_{\max}$. *(Verified at a stable interior equilibrium: $\Delta_{\min} = -6.43 \le \Delta = -2.17 \le \Delta_{\max} = +3.23$.)* □

Theorem 5 (Least-favorable conditional shape)**.** *Fix the equilibrium capacities and q°. Among conditional laws whose density is non-increasing on the FCL zone $(-K_R^*, K_F^* - K_M^*)$ and non-decreasing on the no-FCL zone $(0, K_F^\circ)$ (the Corollary 2 shape class $\mathcal{U}$), the* uniform *(linear CDF) minimizes the carbon saving Δ of (21).*

Proof. By (21), at fixed capacities Δ is *decreasing* in μ_{low}^* (coefficient $-C_M/(R-c)$) and *increasing* in μ_{mid}° (coefficient $\varphi^\circ - q^\circ > 0$), so it is minimized by the *largest* feasible μ_{low}^* and the *smallest*

feasible μ°_{mid}. A non-increasing density on the FCL zone has conditional mean at *most* the interval midpoint, so the largest feasible μ^{*}_{low} over $\mathcal{U}$ is the midpoint; a non-decreasing density on the no-FCL zone has conditional mean at *least* the midpoint, so the smallest feasible μ°_{mid} is the midpoint. Both midpoints are attained *simultaneously* by the uniform (constant density), which therefore minimizes Δ. *(Verified: over* 3000 *random monotone densities the conditional means respect the midpoint bounds in* 3000/3000*, and a continuous tilt from uniform into $\mathcal{U}$ raises Δ — uniform is the minimizer.)* □

Remark 7 (Scope: this is a fixed-capacity statement). Theorem 5 does *not* establish that the uniform is the worst case for the realized saving $\Delta(F)$ across all $F \in \mathcal{U}$, because the equilibrium capacities $K^{*}_{F}, K^{*}_{M}, K^{*}_{R}, K^{\circ}_{F}$ and q° co-vary with the shape of F (Lemma 2 pins only the cutoff *probabilities*). The direction of the surviving capacity effect—the renewable shift—is signed by Theorem 1 ($K^{*}_{R} > K^{\circ}_{R}$), but a full cross-$F$ minimization would have to control the capacity covariation, which we do not close. We therefore state the genuinely distribution-robust certificate only in its honest form: the closed-form $C_R \le C^{\dagger}_{R}$ guarantees $\Delta \ge 0$ for the uniform/power-law family (Proposition 3), and Theorem 5 shows the uniform is the least-favorable *conditional shape*—not yet the least-favorable *distribution*—within $\mathcal{U}$. (Note the uniform residual demand is itself the knife-edge of Proposition 1, with a corner rather than interior K^{*}_{R}; closing the cross-F statement is left open.)

Remark 8 (What this buys). At fixed equilibrium capacities, Theorems 3–5 reduce the carbon effect to a two-mean affine form, bracket it by bounds a regulator can compute from cutoff data, and identify the uniform as the least-favorable conditional shape—so the distributional detail of F controls the *magnitude* of the carbon change at given capacities. Turning the fixed-capacity least-favorable result into a genuine distribution-robust certificate over all of $\mathcal{U}$ requires signing the capacity covariation (Remark 7); that step is open.

6. Reconciling the regimes

The analysis lives in the large-FCL interior regime. Three other regimes appear in the paper: small-FCL ($K^{*}_{M} \le K^{*}_{R}$, Appendix G, with the widened threshold $\bar{C}^{b}_{M} = \bar{C}_{M} + c\,F(0; K^{\circ}_{R})$, paper eq. 62), full-coverage ($R \ge c + C_F + C_M$, zero blackout), and no adoption. They fit into one ordered map.

Proposition 4 (Regime map and unified triple win)**.** *The relevant C_M-thresholds are strictly ordered,*

$$\underbrace{R - c - C_F}_{C^{\text{fc}}_{M}} \;<\; \underbrace{(R-c)\varphi^{\circ}}_{\bar{C}_{M}} \;<\; \underbrace{\bar{C}_{M} + c\,F(0; K^{\circ}_{R})}_{\bar{C}^{b}_{M}}, \tag{23}$$

and as C_M falls the optimum passes through: (0) *no adoption, $C_M \ge \bar{C}^{b}_{M}$;* (I) *small-FCL adoption, $\bar{C}_M \le C_M < \bar{C}^{b}_{M}$;* (II) *large-FCL interior adoption, $C^{\text{fc}}_{M} < C_M < \bar{C}_M$;* (III) *full coverage, $C_M \le C^{\text{fc}}_{M}$ (blackout $\equiv 0$). The profit and blackout parts of the triple win hold in the* large-FCL *regions* (II)–(III) *for every F, since there the cutoff probabilities are distribution-free (Lemma 2); in the small-FCL region* (I) *the cutoffs are not pinned by the FOCs alone (Lemma 2), and the profit/blackout comparison there relies on the Appendix-G analysis (Role 4). The carbon part holds under Condition 1 in* (II)*; in* (III) *it holds trivially, fossil serving only consumers with FCL on renewable surplus. We do* not *claim it "a fortiori" in* (I)*: a smaller program shrinks both sides of (C), and signing the net change there is left to the small-FCL analysis.*

Proof. $\bar{C}_M < \bar{C}^{b}_{M}$ is immediate since $c\,F(0; K^{\circ}_{R}) > 0$. For $C^{\text{fc}}_{M} < \bar{C}_M$,

$$\bar{C}_M - C^{\text{fc}}_{M} = (R-c)\varphi^{\circ} - (R - c - C_F) = C_F - (R-c)\frac{C_F}{\beta + p - c} = C_F\,\frac{\beta + p - R}{\beta + p - c} > 0$$

by Assumption 1 ($\beta + p > R$). Region (III) is the degenerate favorable case of §2: $R \geq c + C_F + C_M$ makes fossil marginal benefit positive up to $\bar{D}$, blackout vanishes, and the carbon comparison reduces to the surplus-absorption argument with no fossil-fed FCL at the margin. The small-FCL utilization effect is bounded by its large-FCL counterpart because $K_M^* \leq K_R^*$ shrinks the fossil-fed dispatch window $(K_F^* - K_M^*, K_F^*)$. □

Equation (23) stitches the regimes: adoption occurs iff $C_M < \bar{C}_M^b$, and the triple-win argument carries across the whole adoption range, strict in (I)–(II) and degenerate-favorable in (III). The remaining small-FCL fine structure (the exact C_M at which $K_M^* = K_R^*$) is the one boundary we leave to Role 4's regime audit.

7. Numerical validation

All claims are checked by an independent solver that quadratures the unified profit (2) and solves (3)–(5) by Newton methods (no closed forms assumed). Highlights:

- **Uniform closed forms.** Carbon eqs. (25)/(26) match quadrature to $\leq 6 \times 10^{-10}$; capacities (8)–(9) satisfy the FOCs to $\leq 10^{-15}$.
- **Knife-edge.** Along the *large-FCL* portion of the re-optimization locus, $\mathrm{MV}(K_R)$ is constant in K_R (verified at $K_R \in \{1, 4, 5\}$ for $\alpha \in \{0.5, 1, 2\}$, e.g. uniform 0.3444 at both $K_R = 1, 4$); the locus leaves this regime for large K_R ($K_M^* < K_R^*$), where MV is no longer constant. Eq. (12) reproduces paper eq. (24).
- **Capacity channel.** On strictly-concave Beta scale-mixtures $\tilde{x} = \tilde{D} - K_R \tilde{S}$ with interior optima, across symmetric, decreasing-density, and right-skewed consumer demand ($D \sim \bar{D} \cdot \mathrm{Beta}(2,2), (1,3), (2,5), (2,8)$), every case has g strictly decreasing, $K_R^* > K_R^\circ$, and the sign of dK_R^*/dC_M matching the shift-ratio bracket of (17). Representative cell ($\mathrm{Beta}(1,3)$): $K_R^* = 3.274 > K_R^\circ = 0.831$, $g(K_F^* - K_M^*) = 0.457 > g(K_F^*) = 0.414$, $dK_R^*/dC_M = -2.99 < 0$ with bracket $-3.1 \times 10^{-4} < 0$. Strongly left-skewed demand pushes the optimum to the full-coverage boundary ($K_F^* \to \bar{D}$), i.e. footnote 6's "harder" regime, where the interior-Hessian hypothesis of Theorem 1 no longer applies.
- **Regime crossing & Hessian.** Scanning C_M from $\bar{C}_M$ downward on $\mathrm{Beta}(2,2)$: the optimum is small-FCL for $C_M \in \{0.77, 0.70, 0.60\}$ ($K_M^* < K_R^*$) and large-FCL for $C_M \in \{0.50, 0.40, 0.30\}$, with $dK_R^*/dC_M < 0$ at *every* point ($-4.28, -2.76, -2.19, -5.21, -21.1, -92.1$; the last is a noisy finite-difference at the knife-edge $K_R^* = 9.2$), confirming the channel across the crossing. The Hessian's three eigenvalues are all strictly negative on the interior path and acquire a zero only as $K_R^* \to \infty$ at the knife-edge ($C_M = 0.30$, $K_R^* = 9.2$ and climbing)—the boundary where Condition 2 lapses and the inequality is anyway immediate.
- **Exactness audit.** `verify_exact.py` checks all 15 analytical formulas (gradients, cross-partials, cofactor, knife-edge MV, distribution-free cutoffs, fossil decomposition) against finite differences of Π: 15/15 agree to 6–16 digits.
- **Viability threshold (Theorem 2).** `verify_threshold.py`: sweeping C_R on $\mathrm{Beta}(2,5)$ right-skewed demand, carbon falls for small C_R (K_R^* unbounded, $\mathbb{E}[\text{fossil}^*] \to 0$, $\Delta = +0.56, +1.07$ at $C_R = 0.25, 0.35$) and rises once C_R crosses $C_R^\dagger \in (0.35, 0.5)$ ($\Delta = -3.68, \ldots, -6.47$ through $C_R = 1.6$). The large-cost anchor (i) is the rigorous, fully general part; the small-cost window is verified here only for scale-mixtures, where two caveats apply—the equilibrium solver is erratic near the knife-edge (we report the clean cold-started points), and for strongly right-skewed demand the no-FCL grid also builds heavy renewables, so the strict fall on $(C_R^\circ, C_R^\dagger)$ can be razor-thin ($\Delta \to 0^+$).
- **Sufficient statistic, bounds, least-favorable (§5).** `verify_carbon_robust.py`: at a stable interior equilibrium the two-moment formula (21) matches direct quadrature to 6×10^{-7}; Δ lies inside the closed-form bounds (22) ($-6.43 \leq -2.17 \leq +3.23$); and over 3000 random

monotone densities the conditional means respect the midpoint bounds in 3000/3000 cases (uniform least-favorable *at fixed capacities*; the cross-distribution version is discussed in §5).

- **Sharpness (algebraic).** The bracket in (17) changes sign with $g(K_F^*) - g(K_F^* - K_M^*)$, so an increasing g would reverse the channel ($K_R^* < K_R^\circ$); the shift-ratio condition is therefore plausibly tight. A realized reversed-channel example proved hard to construct—even at demand–supply correlation $+0.95$ the equilibrium window stayed in the decreasing-g region—so an explicit counterexample is left to Role 3 (comparative statics).

8. Status of the open points

Of the four issues the uniform-only treatment leaves open, three are closed and the fourth is reduced to one boundary; the companion memo (`gap_memo.md`) carries the same accounting.

(1) **Submodularity — closed.** By (15), Π is submodular in (K_R, K_F) with increasing differences in each of $(K_R; K_M)$ and $(K_F; K_M)$; no relabeling fixes both signs, so a lattice argument does not apply. Remark 3 shows this is not an obstruction: Theorem 1's homotopy-plus-cofactor proof is complete under Condition 2 ($\mathcal{H} \prec 0$, verified by eigenvalue computation along the path in §7; it lapses only at the knife-edge, where $K_R^* = \infty > K_R^\circ$ is immediate).

(2) **Endogenous-capacity feedback — closed at fixed capacities.** Lemma 2 pins the comparison *probabilities* distribution-free, and Theorem 1 signs the one surviving level effect ($K_R^* > K_R^\circ$); at fixed capacities (C) is then a complete affine representation (Theorem 3, Remark 5). The residual cross-distribution capacity covariation (Remark 7) is signed in direction but not fully closed.

(3) $C_R^\dagger$ **— general upper anchor, explicit window.** Theorem 2 proves $\{C_R : \Delta \geq 0\}$ is bounded above for *every* F (the rigorous large-cost limit $\Delta \to -U_0 < 0$), and exhibits a nonempty strict-fall window $(C_R^\circ, C_R^\dagger)$ for the knife-edge / power-law family, with explicit $C_R^\dagger = \mathrm{MV}$, $C_R^\circ = C_F^2/2(\beta + p - c)$ (Proposition 3). *This upgrades the earlier heuristic "the two curves cross" to a proof of the upper anchor; a general-F nonempty window remains conjectural.*

(4) **Boundary regimes — ordered, one crossover open.** Proposition 4 proves the strict ordering $C_M^{\mathrm{fc}} < \bar{C}_M < \bar{C}_M^b$ and the triple win in the large-FCL and full-coverage regions; the small-FCL region (I) and the exact crossover $K_M^* = K_R^*$ rest on the Appendix-G analysis, handed to Role 4.